

R Coding Sample (Royals Project)

Nick Skiera

2023-11-14

Querying Prompts

Using the attached dataset of throws to first base, please answer the following questions. Please utilize SQL, R or Python to aggregate the data. Attach all code and visualizations that you used throughout your entire process.

```
dat <- read.csv("~/Downloads/dataset_2024.csv")
head(dat)
```

```
##   throw_id team_id fielder_id fielder_position thrower_id thrower_position
## 1         3      11       400                6         400                6
## 2         6      11       228                5         390                4
## 3         7       8       415                4         415                4
## 4         8       8       308                1         308                1
## 5        10       1       314                4         300                6
## 6        11       2       312                4         312                4
##   receiver_id receiver_position exchange_time throw_pos_x throw_pos_y
## 1          63                3         1.533 -60.116123 132.18728
## 2          63                3         0.534 -0.562563 122.82194
## 3         143                3         1.266  1.598751 123.03693
## 4         143                3         1.800 25.403185  59.92459
## 5         514                3         0.733  8.957441 126.16588
## 6         695                3         1.200 15.075073 138.72246
##   throw_velo_x throw_velo_y throw_velo_z batter_pos_x_at_throw
## 1    60.04756  -33.596364    8.583470         26.91291
## 2    51.29392  -45.316519    8.831870         43.39641
## 3    50.04965  -41.717447    5.986785         25.30078
## 4    24.01967   3.289158   12.783662         19.69922
## 5    51.63067  -63.361730    7.811719         40.11830
## 6    41.26329  -62.250650    6.230779         39.70496
##   batter_pos_y_at_throw batter_velo_at_throw bounce_pos_x bounce_pos_y
## 1         26.55610         25.78753      38.06139      69.47463
## 2         44.29049         28.64788           NA           NA
## 3         29.97072         25.13043           NA           NA
## 4         15.24991         18.50923           NA           NA
## 5         41.43340         27.31851           NA           NA
## 6         41.80575         27.57804           NA           NA
##   bounce_velo_x bounce_velo_y bounce_velo_z receiver_pos_x receiver_pos_y
## 1    32.81451    -25.4525    7.330194     56.19932     60.18778
## 2         NA         NA         NA     59.01313     65.65794
## 3         NA         NA         NA     61.03749     64.42913
## 4         NA         NA         NA     63.14208     65.64872
## 5         NA         NA         NA     59.77040     64.04606
```

```
## 6      NA      NA      NA      59.62286      63.94086
## receiver_dist_from_1b throw_deflected_by_receiver start_state end_state
## 1      8.202015      0      ____1      1____1
## 2      5.047566      0      1____1      1____2
## 3      2.719261      0      123_1      ____3
## 4      2.069797      0      ____1      ____2
## 5      3.890504      0      123_1      ____3
## 6      4.028031      0      ____0      ____1
## runs_on_play batter_result
## 1      0      first
## 2      0      first
## 3      0      out
## 4      0      out
## 5      0      out
## 6      0      out
```

1. Which 5 infielders had the quickest exchange times on throws to first base?

```
infielders <- c(2,4,5,6)
dat %>%
  filter(thrower_position %in% infielders) %>% #filtering out outfielders and pitchers
  filter(exchange_time > 0) %>%
  dplyr::select(thrower_id, thrower_position, exchange_time) %>%
  arrange(exchange_time) %>%
  top_n(-5)
```

Selecting by exchange_time

```
## thrower_id thrower_position exchange_time
## 1      592      6      0.034
## 2      396      5      0.067
## 3      112      4      0.167
## 4      159      4      0.167
## 5      85      4      0.200
## 6      687      5      0.200
## 7      267      4      0.200
## 8      190      2      0.200
```

The 5 infielders that are not pitchers with the quickest exchange times are throwers 592, 396, 112/159 and 85/687/267/190. I removed pitchers because they are not position players we typically evaluate when looking at fielding. I removed exchange times that are 0 because those are likely glove flips or barehanded players.

```
infielders_w_pitchers <- c(1,2,4,5,6)
dat %>%
  filter(thrower_position %in% infielders_w_pitchers) %>% #filtering out only outfielders
  filter(exchange_time > 0) %>%
  dplyr::select(thrower_id, thrower_position, exchange_time) %>%
  arrange(exchange_time) %>%
  top_n(-5)
```

Selecting by exchange_time

```
## thrower_id thrower_position exchange_time
## 1      1      1      0.033
## 2      592      6      0.034
## 3      292      1      0.034
## 4      658      1      0.067
```

```
## 5          396          5          0.067
```

However, if we do include pitchers then the infielders with the quickest exchange times on throws to first are throwers 1, 592, 292, 658 and 396

2. The infield coach wants to see which teams made the most errant throws to first base. An errant throw is described as a throw that bounced and resulted in the runner being safe. Please create a basic visual that you would present to the infield coach to present your findings.

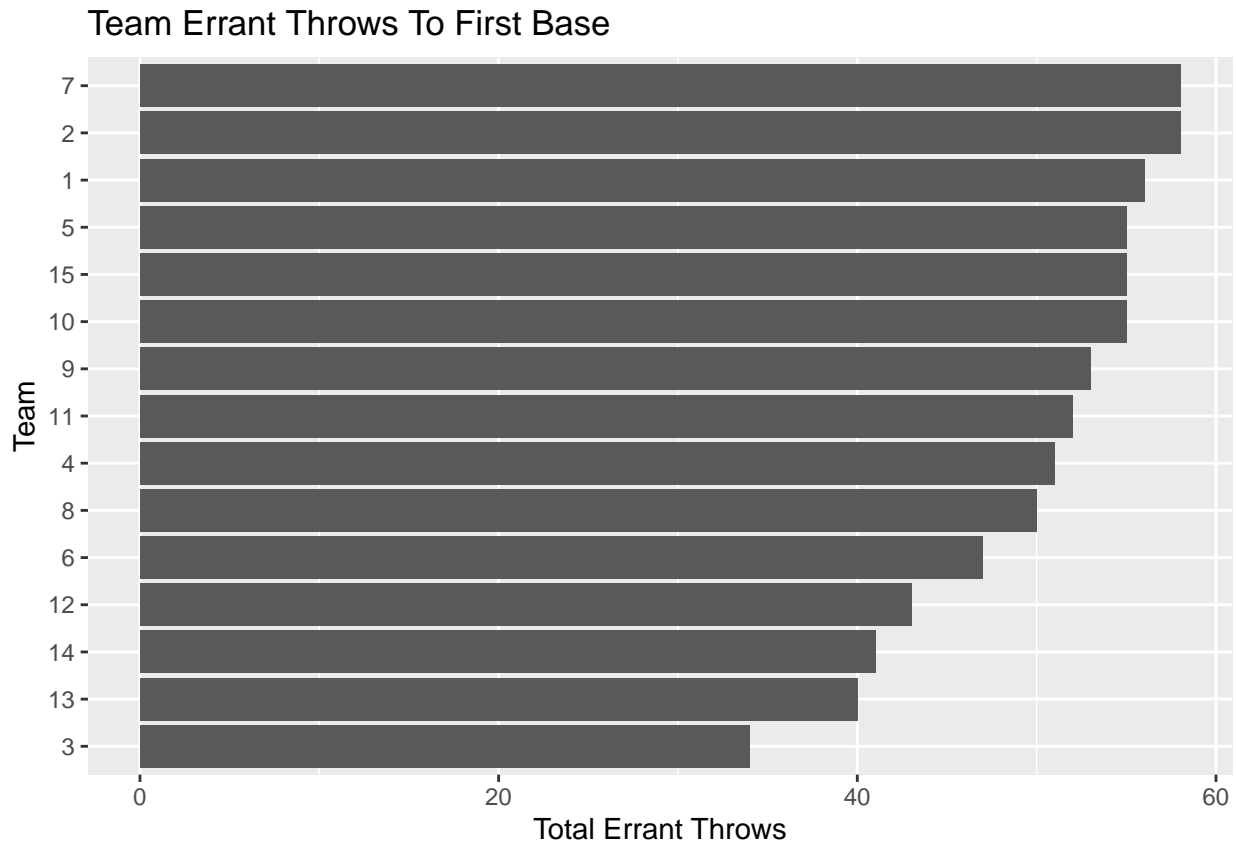
```
err_throw <- dat %>%
  filter(thrower_position %in% infielders_w_pitchers) %>%
  mutate(errant_throw = as.factor(ifelse(!is.na(bounce_pos_x) & batter_result != "out", 1, 0))
  filter(errant_throw == 1) %>% #will remove all throws that don't bounce
  arrange(team_id)

team_err_throws <- c()
for(i in 1:15) {
  team_err_throws[i] <- sum(err_throw$team_id == i)
}

head(df_err <- data.frame(
  team_id = as.character(unique(err_throw$team_id)),
  team_err_throws = team_err_throws) %>%
  mutate(team_order = fct_reorder(team_id, team_err_throws)))

##   team_id team_err_throws team_order
## 1      1          56          1
## 2      2          58          2
## 3      3          34          3
## 4      4          51          4
## 5      5          55          5
## 6      6          47          6

ggplot(df_err, aes(x = team_order, y = team_err_throws)) +
  geom_bar(stat = "identity") +
  labs(title = "Team Errant Throws To First Base") +
  xlab("Team") +
  ylab("Total Errant Throws") +
  coord_flip()
```



3. Looking at all infield throws to first base, given that the distance of the throw to first base was in the top 90th percentile, what team had the best average exchange time? Which team had the largest variation in exchange time on these throws?

```
dist_throw_1b <- dat %>%
  filter(thrower_position %in% infielders_w_pitchers) %>%
  filter(exchange_time > 0) %>%
  mutate(distance_of_throw = sqrt((receiver_pos_x - throw_pos_x)^2 + (receiver_pos_y - throw_pos_y)^2))
  mutate(percentile_rank = ntile(desc(distance_of_throw), 10)) %>%
  arrange(desc(distance_of_throw))

head(throws_90_pctile <- dist_throw_1b %>%
  filter(percentile_rank == 1))
```

##	throw_id	team_id	fielder_id	fielder_position	thrower_id	thrower_position
## 1	1313	5	112	4	112	4
## 2	2778	5	446	4	446	4
## 3	18712	6	566	6	566	6
## 4	15975	15	426	5	426	5
## 5	27885	15	159	4	159	4
## 6	25832	7	673	4	673	4

##	receiver_id	receiver_position	exchange_time	throw_pos_x	throw_pos_y
## 1	766	3	1.767	-104.49645	403.5053
## 2	766	3	0.833	91.56045	403.6370
## 3	25	3	1.734	132.01699	375.0706
## 4	568	3	0.666	169.96265	329.2181
## 5	313	3	0.866	214.40527	283.8712
## 6	615	3	2.267	223.92484	272.7766

```
##   throw_velo_x throw_velo_y throw_velo_z batter_pos_x_at_throw
## 1   -24.44178   -25.40783    6.780555      34.75114
## 2   -23.14839    16.59581   11.931548      39.28094
## 3   -20.50401    25.14487    3.784122      33.14163
## 4   -24.38325    19.03692   -4.884555      19.41157
## 5   -23.56285    27.72982    6.858485      34.98759
## 6   -27.55181    34.07074    8.871604      31.61921
##   batter_pos_y_at_throw batter_velo_at_throw bounce_pos_x bounce_pos_y
## 1           34.31264           26.31457           NA           NA
## 2           38.94313           28.01125           NA           NA
## 3           34.08661           25.23992           NA           NA
## 4           20.03403           23.30695           NA           NA
## 5           33.14406           24.66136           NA           NA
## 6           27.91198           19.46385           NA           NA
##   bounce_velo_x bounce_velo_y bounce_velo_z receiver_pos_x receiver_pos_y
## 1           NA           NA           NA      60.32098      63.92279
## 2           NA           NA           NA      60.16601      64.60350
## 3           NA           NA           NA      60.44282      63.03786
## 4           NA           NA           NA      59.42645      63.54653
## 5           NA           NA           NA      61.28857      63.59752
## 6           NA           NA           NA      62.40411      64.56134
##   receiver_dist_from_1b throw_deflected_by_receiver start_state end_state
## 1           3.330693           0           1___1      ___3
## 2           3.604850           0           ___2      ___3
## 3           3.252932           0           _2__2      ___3
## 4           4.214187           0           ___0      ___1
## 5           2.351422           0           ___0      ___1
## 6           1.541444           0           ___0      ___1
##   runs_on_play batter_result distance_of_throw percentile_rank
## 1           0           out           377.4667           1
## 2           0           out           340.4839           1
## 3           0           out           320.1363           1
## 4           0           out           287.7492           1
## 5           0           out           268.2634           1
## 6           0           out           263.5195           1
```

```
avg_exch_time <- c()
for (i in 1:15) {
  avg_exch_time[i] <- mean(throws_90_pctile$exchange_time[which(throws_90_pctile$team_id == i)])
}
var_exch_time <- c()
for (i in 1:15) {
  var_exch_time[i] <- var(throws_90_pctile$exchange_time[which(throws_90_pctile$team_id == i)])
}
```

```
(df_throw_dist <- data.frame(
  team_id = as.character(unique(throws_90_pctile$team_id)),
  avg_exch_time = avg_exch_time,
  var_exch_time = var_exch_time) %>%
  arrange(avg_exch_time))
```

```
##   team_id avg_exch_time var_exch_time
## 1       2    1.132352    0.1388616
## 2      14    1.157395    0.1198779
## 3       6    1.161526    0.1090549
```

```
## 4      5      1.164255      0.1383404
## 5      4      1.172265      0.1290435
## 6     13      1.185328      0.1142765
## 7     10      1.193778      0.1687950
## 8      7      1.194088      0.1626194
## 9     12      1.202600      0.1338400
## 10     3      1.212613      0.1256264
## 11    15      1.215025      0.1348278
## 12     8      1.228241      0.1462788
## 13     9      1.241426      0.1387692
## 14    11      1.255300      0.1477178
## 15     1      1.322391      0.1051598
```

```
df_throw_dist$team_id[which(df_throw_dist$avg_exch_time == min(df_throw_dist$avg_exch_time))]
```

```
## [1] "2"
```

```
df_throw_dist$team_id[which(df_throw_dist$var_exch_time == max(df_throw_dist$var_exch_time))]
```

```
## [1] "10"
```

Team 2 had the lowest average exchange time at 1.132352 and Team 10 had the largest variation in exchange time at 0.1687950.

- Given that a throw was made less than 100 feet from first base, is there a correlation between throw velocity and throw distance? Provide a basic visual alongside a brief explanation.

```
head(short_throw <- dist_throw_1b %>%
  filter(distance_of_throw < 100) %>%
  mutate(throw_velo = sqrt(throw_velo_x^2 + throw_velo_y^2 + throw_velo_z^2)))
```

```
##   throw_id team_id fielder_id fielder_position thrower_id thrower_position
## 1    14638      8        26              5         26              5
## 2    24105      7       757              6       757              6
## 3    27977      4       379              6       379              6
## 4    18143      3       653              5       653              5
## 5    24302      7       757              6       757              6
## 6     9638      3       653              5       653              5
##   receiver_id receiver_position exchange_time throw_pos_x throw_pos_y
## 1         863              3         0.933   -37.95813   73.84991
## 2         820              3         0.900   -32.63508   97.67744
## 3         397              3         0.633   -20.44641  120.65644
## 4         103              3         1.133   -38.61237   75.88515
## 5         339              3         0.634   -20.72630  121.82080
## 6         474              3         1.300   -32.62485   99.70371
##   throw_velo_x throw_velo_y throw_velo_z batter_pos_x_at_throw
## 1    77.36064   -3.579660    2.471836          36.79104
## 2    72.93405  -18.626071    9.118424          34.14535
## 3    54.93388  -35.254666   10.689079          26.00501
## 4    79.23197   -5.978675    6.226513          34.69233
## 5    61.85791  -39.462358    9.019587          36.83248
## 6    67.02635  -22.415908    7.923176          33.04008
##   batter_pos_y_at_throw batter_velo_at_throw bounce_pos_x bounce_pos_y
## 1          38.84842          28.42306      43.76895      66.43157
## 2          36.28389          28.56228           NA           NA
## 3          27.19864          22.19235           NA           NA
## 4          35.19163          25.34513           NA           NA
```

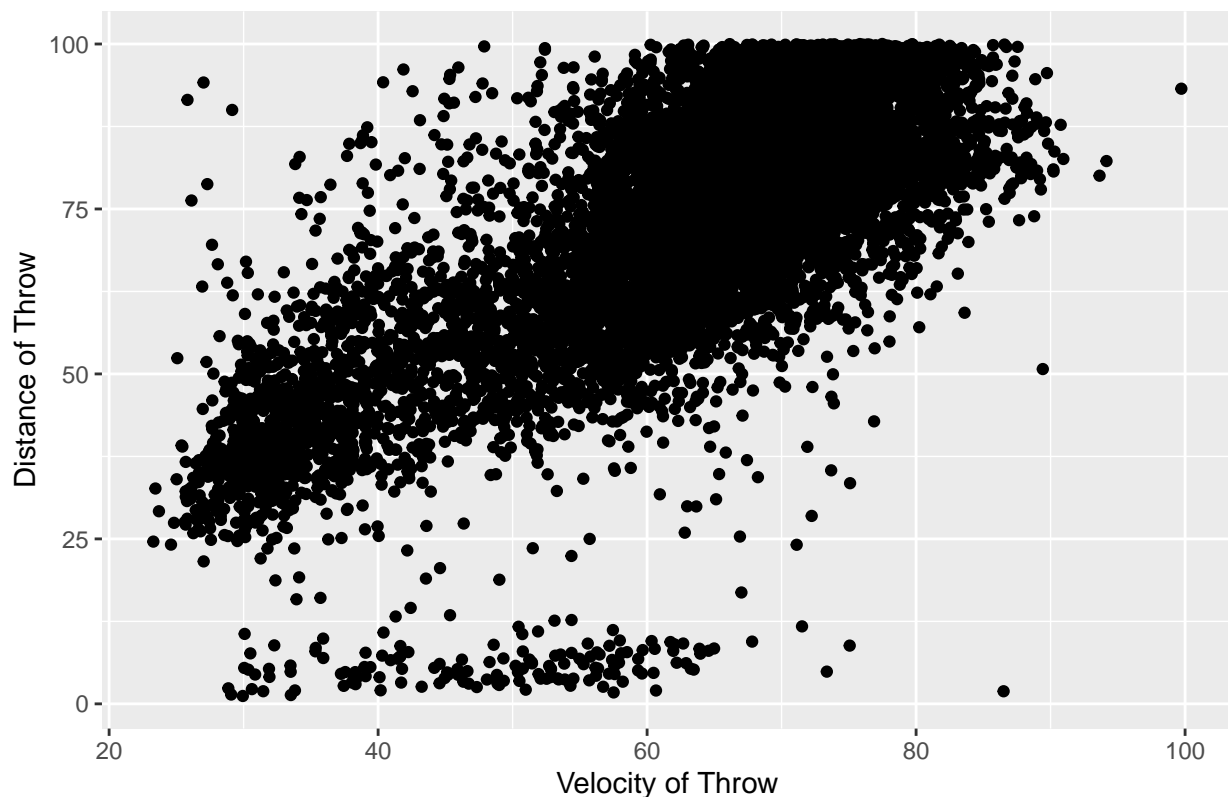
```
## 5          37.74694          25.23409          NA          NA
## 6          33.88707          24.75661          NA          NA
##  bounce_velo_x bounce_velo_y bounce_velo_z receiver_pos_x receiver_pos_y
## 1          52.96815        -8.404934         6.527076         61.63916         64.89060
## 2           NA           NA           NA         62.49966         66.90812
## 3           NA           NA           NA         61.90322         63.96179
## 4           NA           NA           NA         60.53937         63.13401
## 5           NA           NA           NA         61.08303         64.39102
## 6           NA           NA           NA         61.72682         66.75446
##  receiver_dist_from_1b throw_deflected_by_receiver start_state end_state
## 1           2.359400              0         ____2         ____3
## 2           3.461598              0         ____0         ____1
## 3           1.766031              0         ____2         ____3
## 4           3.141200              0         ____2         ____3
## 5           2.664716              0         ____1         ____2
## 6           3.655276              0         _2__2         ____3
##  runs_on_play batter_result distance_of_throw percentile_rank throw_velo
## 1           0           out           99.99945              5       77.48285
## 2           0           out           99.98685              5       75.82514
## 3           0           out           99.97871              5       66.14286
## 4           0           out           99.96829              5       79.70081
## 5           0           out           99.95472              5       73.92585
## 6           0           out           99.93943              5       71.11808

cor(short_throw$throw_velo, short_throw$distance_of_throw, use = "complete.obs")

## [1] 0.7380965

ggplot(short_throw, aes(x = throw_velo, y = distance_of_throw)) +
  geom_point() +
  labs(title = "Scatter Plot of Throw Distance on Throw Velocity") +
  xlab("Velocity of Throw") +
  ylab("Distance of Throw")
```

Scatter Plot of Throw Distance on Throw Velocity



We can see based on the plot that, outside of a few outliers, there is clearly a strong positive relationship between the distance of the throw and the velocity of the throw with a correlation of 0.7380965.

Modeling Project

While often routine, an infielder making a timely and accurate throw to first base is a skill that is critical to the outcome of a game. Arm strength, exchange time, velocity, and the first baseman's ability to receive an errant throw all determine whether or not an out is made on the play. We have attached a dataset of throws to first base for you to analyze. We would like to see you build some sort of model based on this data that evaluates the talent of a subset of infielders. We will be evaluating your submission on four components of your output:

1. Your modeling approach and creativity.
2. Your model evaluation process.
3. Your ability to perform a skill assessment of some group of players involved.
4. Your creation of a player-evaluation tool or other presentation layer that would be presented to a less technical audience.

If there is additional information that you think would clarify this problem or strategies you would implement if you had more time, please detail that as well.

This project should take approximately 8-12 hours to complete. Please include all code written (in either R or Python) for this analysis, as well as your final product.

```
new_dat <- dat %>%
  filter(thrower_position %in% infielders) %>% #taking out outfielders and pitchers
  subset(thrower_position == fielder_position) %>% #isolates data to just fielders who made throws to 1
  mutate(thrower_err_throw = as.factor(ifelse(!is.na(bounce_pos_x), 1, 0))) %>% #variable for thrower b
  mutate(first_base_save = as.factor(ifelse(!is.na(bounce_pos_x) & batter_result == "out", 1, 0))) %>%
  mutate(throw_velo = sqrt(throw_velo_x^2 + throw_velo_y^2 + throw_velo_z^2)) %>% #create overall throw
```



```

mutate(num_outs = as.numeric(substr(start_state, 5, 5))) %>% #create variable for number of outs
mutate(runner_on_third = as.factor(ifelse(substr(start_state, 3, 3) == "3", 1, 0))) %>% #create varia
mutate(distance_of_throw = sqrt((receiver_pos_x - throw_pos_x)^2 + (receiver_pos_y - throw_pos_y)^2))
mutate(batter_result = as.factor(ifelse(batter_result == "out", 1, 0))) %>% #changes batter_result to
remove_rownames() %>% column_to_rownames( var = "throw_id") %>%
dplyr::select(-c(bounce_pos_x, bounce_pos_y, bounce_velo_x, bounce_velo_y, bounce_velo_z, throw_velo_
new_dat <- na.omit(new_dat)

head(new_dat)

```

```

##      team_id thrower_id thrower_position receiver_id exchange_time
## 3         11        400                6         63         1.533
## 7          8        415                4        143         1.266
## 11         2        312                4        695         1.200
## 14         2        396                5        695         0.966
## 15         1        314                4        514         0.967
## 16         1        201                5        514         1.633
##      batter_pos_x_at_throw batter_pos_y_at_throw batter_velo_at_throw
## 3                26.91291                26.55610                25.78753
## 7                25.30078                29.97072                25.13043
## 11               39.70496                41.80575                27.57804
## 14               36.31824                39.76688                29.64266
## 15               38.21505                38.47144                26.57897
## 16               27.59490                29.82228                24.41569
##      receiver_dist_from_1b throw_deflected_by_receiver runs_on_play batter_result
## 3                8.202015                0                0                0
## 7                2.719261                0                0                1
## 11               4.028031                0                0                1
## 14               4.419242                0                0                1
## 15               2.929510                0                0                1
## 16               11.914992                0                1                0
##      thrower_err_throw first_base_save throw_velo num_outs runner_on_third
## 3                1                0    69.34047         1                0
## 7                0                0    65.43053         1                1
## 11               0                0    74.94414         0                0
## 14               0                0    70.72622         0                0
## 15               0                0    69.82297         0                0
## 16               1                0    75.34457         1                1
##      distance_of_throw
## 3                136.79624
## 7                 83.47357
## 11                87.04478
## 14                70.42454
## 15                66.81671
## 16                140.20326

```

```
dim(new_dat)
```

```
## [1] 15385    18
```

I removed outfielders and pitchers because we are evaluating infielders and the main job of pitchers is not fielding baseballs and making good throws to first base. I also removed plays where the thrower didn't field the ball so we can eliminate abnormal plays such as ricochets and double plays involving multiple infielders other than the receiver. I created the variables thrower errant throw, first base save, throw velocity, number of outs, runner on third and distance of throw because I believe that this will add a complexity while also

reducing the amount of variables in the data set. I created thrower errant throw because I wanted to eliminate the bounce coordinate variables because they contain NAs. First base save shows us if the first basemen saved the errant throw. I used the velocity coordinates on the throw to create an overall throwing velocity variable which will reduce the dimensionality while maintaining the value of the variables. I wanted to remove the variables start_state and end_state because they are character variables so I extracted the data I deemed important which was number of outs and if there was a runner on third. The number of outs and having a runner on third can add pressure to make a good throw depending on the situation. I also wanted to remove the thrower and receiver coordinates but to maintain their importance I created the distance of throw variable. Finally, I mutated the variable batter_result to be a response factor of whether the batter was thrown out/success (1) or reached base/failure (0).


```
set.seed(100)

#splitting data into training and test data for cross validation
train_inx <- sample(seq_len(nrow(new_dat)), size = 0.7*nrow(new_dat))
dat_train <- new_dat[train_inx, -c(1, 2, 4)] #removing player ids manually
dat_test <- new_dat[-train_inx, -c(1, 2, 4)]
```

I took out all identification variables (team_id, thrower_id and receiver_id) because these are not predictors. We have 14 predictors. I split the data set into training and test data to use in cross validation and to compare models in-sample and out-of-sample classification accuracy. I hope to find a model with a high classification accuracy in both to have a good fit. A model with a high in-sample classification accuracy and low out-of-sample classification accuracy is over fit because it conforms to the data it was built upon well but not to new data. The opposite scenario means that the model would be under fit because it is able to conform to any data but not the data it was built upon. This is similar to bias-variance trade-off.

```
set.seed(100)

cv_lasso <- cv.glmnet(y = dat_train$batter_result, x = as.matrix(dat_train[, -9]), alpha = 1, nfolds = 10)
autoplot(cv_lasso)
```



```
lasso_1se <- glmnet(y = dat_train$batter_result, x = as.matrix(dat_train[,-9]), alpha = 1, lambda = cv_1se_lambda)
coef(lasso_1se)
```

```
## 15 x 1 sparse Matrix of class "dgCMatrix"
##                               s0
## (Intercept)                   27.29658281
## thrower_position               -0.08079468
## exchange_time                  0.75377407
## batter_pos_x_at_throw          -0.18328416
## batter_pos_y_at_throw          -0.11338428
## batter_velo_at_throw           -0.43075000
## receiver_dist_from_1b          -0.07178555
## throw_deflected_by_receiver   -6.70331836
## runs_on_play                   -2.13608634
## thrower_err_throw              -5.39103958
## first_base_save                 9.15897327
## throw_velo                     0.07547682
## num_outs                       .
## runner_on_third                 0.77420350
## distance_of_throw              -0.07449516
```

```
pred_lasso <- predict(lasso_1se, newx = as.matrix(dat_train[,-9]), type = "class")
mean(pred_lasso == dat_train$batter_result)
```

```
## [1] 0.9700994
```

I used cross validation in lasso regression with the desire for dimension reduction and see that it actually eliminated the variable num_outs from the model. I will take num_outs out of the data frame for the rest of

the tested models. We now have 13 predictors

```
dat_train1 <- dat_train[,-13]
dat_test1 <- dat_test[,-13]
```

```
mod_logit <- glm(batter_result ~ ., data = dat_train1, family = "binomial")
summary(mod_logit)
```

```
##
## Call:
## glm(formula = batter_result ~ ., family = "binomial", data = dat_train1)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    35.509032    1.614840   21.989 < 2e-16 ***
## thrower_position -0.175949    0.079499   -2.213 0.026882 *
## exchange_time     0.909231    0.157887    5.759 8.47e-09 ***
## batter_pos_x_at_throw -0.239710    0.019398  -12.357 < 2e-16 ***
## batter_pos_y_at_throw -0.150832    0.019687   -7.661 1.84e-14 ***
## batter_velo_at_throw  -0.569881    0.038582  -14.771 < 2e-16 ***
## receiver_dist_from_1b -0.089965    0.019775   -4.549 5.38e-06 ***
## throw_deflected_by_receiver -9.584015    0.877126  -10.927 < 2e-16 ***
## runs_on_play        -3.743544    0.623854   -6.001 1.97e-09 ***
## thrower_err_throw1  -20.946685   382.564552   -0.055 0.956335
## first_base_save1     44.595619   574.973767    0.078 0.938177
## throw_velo           0.110000    0.009185   11.976 < 2e-16 ***
## runner_on_third1     2.211648    0.596827    3.706 0.000211 ***
## distance_of_throw    -0.102763    0.005325  -19.300 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 6185.5  on 10768  degrees of freedom
## Residual deviance: 1799.7  on 10755  degrees of freedom
## AIC: 1827.7
##
## Number of Fisher Scoring iterations: 18
```

```
pred_logit <- predict(mod_logit, type = "response") > 0.5
mean(pred_logit == as.numeric(dat_train1$batter_result))
```

```
## [1] 0.0218219
```

```
out_pred_logit <- predict(mod_logit, newdata = dat_test1, type = "response") > 0.5
mean(out_pred_logit == as.numeric(dat_test1$batter_result))
```

```
## [1] 0.02058059
```

We can see that a logistic model is not great with an in-sample accuracy of 0.0218219 and out-of-sample accuracy of 0.02058059. This makes sense as this model is not cross validated. We will move forward to a K-Nearest-Neighbors model using cross validation.

```
set.seed(100)
```

```
trctrl <- trainControl(method = "cv", number = 10)
knn_fit <- train(batter_result ~ ., data = dat_train1, method = "knn", trControl = trctrl, tuneGrid = e
```

```
knn_fit$bestTune
```

```
## k  
## 7 7
```

```
mod_knn <- knn(train = dat_train1[, -9], cl = dat_train1$batter_result, test = dat_train1[, -9], k = knn_
```

```
mean(mod_knn == dat_train1$batter_result)
```

```
## [1] 0.9605349
```

```
mean(mod_knn == dat_test1$batter_result)
```

```
## [1] 0.8704615
```

This prediction is a lot better with an in-sample accuracy of 0.9605349 and out-of-sample accuracy of 0.8704615. This model will work but we will test an LDA model to see if we can still improve.

```
mod_lda <- lda(batter_result ~., data = dat_train1, prior = rep(1, 2)/2)  
pred_lda <- predict(mod_lda, dat_train1)
```

```
mean(pred_lda$class == dat_train1$batter_result)
```

```
## [1] 0.9594206
```

```
mean(pred_lda$class == dat_test1$batter_result)
```

```
## [1] 0.8699044
```

LDA predicts worse than KNN but still great with an in-sample accuracy of 0.9594206 and out-of-sample accuracy of 0.8699044. Now I will combine all of these models on an in-sample ROC plot and calculate the Areas Under the Curves.

```
set.seed(100)
```

```
prob_logit <- predict(mod_logit, type = "response")  
prob_lasso <- predict(lasso_1se, newx = as.matrix(dat_train[, -9]), type = "response")  
prob_knn <- 1 - attributes(knn(train = dat_train1[, -9], cl = dat_train1$batter_result, test = dat_train
```

```
prob_lda <- predict(mod_lda, dat_train1)$posterior[, 2]  
df_roc <- data.frame(logit = prob_logit,  
                     lasso = prob_lasso,  
                     knn = prob_knn,  
                     lda = prob_lda,  
                     batter_result = dat_train$batter_result)  
names(df_roc) <- c("logit", "lasso", "knn", "lda", "batter_result" )  
rocobj <- roc(batter_result ~ logit + lasso + knn + lda, data = df_roc)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```


```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls > cases
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
ggroc(rocobj)
```



```
df_auc <- data.frame(logit = auc(rocobj$logit),
                     lasso = auc(rocobj$lasso),
                     knn = auc(rocobj$knn),
                     lda = auc(rocobj$lda))
```

```
df_auc
```

```
##      logit      lasso      knn      lda
## 1 0.9792219 0.9776247 0.8193673 0.9768522
```


The logistic model actually has the highest AUC with a 0.9792219. However, due to the low accuracies, I will avoid this model. This might be due to the nature of ROC and how it assigns scores. However, LDA has a great AUC with a 0.9768522 which pairs well with its classification accuracy.

```
prob_logit_test <- predict(mod_logit, newdata = dat_test1, type = "response")
prob_lass_test <- predict(lasso_1se, newx = as.matrix(dat_test[, -9]), type = "response")
prob_knn_test <- 1 - attributes(knn(train = dat_train1[, -9], cl = dat_train1$batter_result, test = dat_test1[, -9]))
prob_lda_test <- predict(mod_lda, dat_test1)$posterior[, 2]
```

```
df_roc_test <- data.frame(logit = prob_logit_test,
                          lasso = prob_lass_test,
                          knn = prob_knn_test,
                          lda = prob_lda_test,
                          Test = dat_test1$batter_result)
names(df_roc_test) <- c("logit", "lasso", "knn", "lda", "batter_result")
rocobj_test <- roc(batter_result ~ logit + lasso + knn + lda, data = df_roc_test)
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
## Setting levels: control = 0, case = 1
## Setting direction: controls > cases
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
```

```
ggroc(rocobj_test)
```



```
df_auc_test <- data.frame(logit = auc(rocobj_test$logit),
                           lasso = auc(rocobj_test$lasso),
                           knn = auc(rocobj_test$knn),
                           lda = auc(rocobj_test$lda))
```

```
df_auc_test
```

```
##      logit      lasso      knn      lda
## 1 0.9738865 0.9723764 0.7172281 0.973346
```

We can also see that LDA performs great with an AUC of 0.973346 that pairs well with its out-of-sample classification accuracy

```
mod_lda
```


```
## Call:
```

```
## lda(batter_result ~ ., data = dat_train1, prior = rep(1, 2)/2)
##
## Prior probabilities of groups:
## 0 1
## 0.5 0.5
##
## Group means:
## thrower_position exchange_time batter_pos_x_at_throw batter_pos_y_at_throw
## 0 5.228031 0.9713026 36.74825 37.93253
## 1 4.978419 1.1905809 32.89720 33.40297
## batter_velo_at_throw receiver_dist_from_1b throw_deflected_by_receiver
## 0 27.01667 7.666191 0.2280311457
## 1 24.75818 3.292706 0.0005065856
## runs_on_play thrower_err_throw1 first_base_save1 throw_velo runner_on_third1
## 0 0.11234705 0.47163515 0.00000000 70.96213 0.08120133
## 1 0.01550152 0.04620061 0.04620061 70.34564 0.07588652
## distance_of_throw
## 0 114.22989
## 1 96.50439
##
## Coefficients of linear discriminants:
## LD1
## thrower_position -0.003170845
## exchange_time 0.343448596
## batter_pos_x_at_throw -0.040479914
## batter_pos_y_at_throw -0.033499128
## batter_velo_at_throw -0.021273202
## receiver_dist_from_1b -0.020249905
## throw_deflected_by_receiver -3.112544926
## runs_on_play -0.918889843
## thrower_err_throw1 -5.027873054
## first_base_save1 5.577849448
## throw_velo 0.027130156
## runner_on_third1 0.233984265
## distance_of_throw -0.018695688
```

We can see based on the coefficients of the model that many of the variables actually have a negative coefficient in predicting the result of the batter with a first basement save, velocity of throw, exchange time and having a runner on third having a positive coefficient.

```
plot(pred_lda$x, rep(0, length(pred_lda$x)), col = dat_train1$batter_result, pch = 16, main = "LDA Model",
legend("topright", legend = levels(dat_train1$batter_result), col = 1:2, pch = 16)
```


LDA Model Plot



Linear Discriminant

We can see in this plot that there is some overlap between the two results but is overall separated pretty well. This means that the model isn't perfect at predicting the result of a throw to first base which makes sense as baseball is an obscure game that can result in plays like this that should result in an out but don't. Plays that don't result in an out classify more with the negative linear discriminant values and plays that do result in an out correspond more with the positive linear discriminant values.